Question 101 (Describe how to work with non-relational data on Azure)

Your company is designing an application that will write a high volume of JSON data and will have an application-defined schema.
Which type of data store should you use?

A. columnar
B. key/value
C. document
D. graph

Answer : B

Explanation:
A key/value store associates each data value with a unique key.
An application can store arbitrary data as a set of values. Any schema information must be provided by the application. The key/value store simply retrieves or stores the value by key.
Reference:
https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview

Question 102 (Describe how to work with non-relational data on Azure)

You need to recommend a non-relational data store that is optimized for storing and retrieving text files, videos, audio streams, and virtual disk images. The data store must store data, some metadata, and a unique ID for each file.
Which type of data store should you recommend?

A. key/value
B. columnar
C. object
D. document

Answer : C

Explanation:
Object storage is optimized for storing and retrieving large binary objects (images, files, video and audio streams, large application data objects and documents, virtual machine disk images). Large data files are also popularly used in this model, for example, delimiter file (CSV), parquet, and ORC. Object stores can manage extremely large amounts of unstructured data.
Reference:
https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview

Question 103 (Describe how to work with non-relational data on Azure)

Your company is designing a data store for internet-connected temperature sensors.
The collected data will be used to analyze temperature trends.
Which type of data store should you use?

A. relational
B. time series
C. graph
D. columnar

Answer : B

Explanation:
Time series data is a set of values organized by time. Time series databases typically collect large amounts of data in real time from a large number of sources.
Updates are rare, and deletes are often done as bulk operations. Although the records written to a time-series database are generally small, there are often a large number of records, and total data size can grow rapidly.
Reference:
https://docs.microsoft.com/en-us/azure/architecture/guide/technology-choices/data-store-overview

Question 104 (Describe how to work with non-relational data on Azure)

HOTSPOT -
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.
Hot Area:

Answer Area

Statements	Yes	No
Copying data to Azure Data Lake Storage from another Azure region results in extra cost.	○	○
You can use blob, table, and file storage in the same Azure Storage account.	○	○
You implement Azure Data Lake Storage by creating an Azure Storage account.	○	○

Answer :

Answer Area

Statements	Yes	No
Copying data to Azure Data Lake Storage from another Azure region results in extra cost.	○	○
You can use blob, table, and file storage in the same Azure Storage account.	○	○
You implement Azure Data Lake Storage by creating an Azure Storage account.	○	○

Reference:
https://docs.microsoft.com/en-us/azure/data-lake-store/data-lake-store-get-started-portal https://docs.microsoft.com/en-us/azure/storage/common/storage-account-overview https://azure.microsoft.com/en-us/pricing/details/bandwidth/

Question 105 (Describe how to work with non-relational data on Azure)

HOTSPOT -
To complete the sentence, select the appropriate option in the answer area.
Hot Area:

Answer Area

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a [▼]
- graph.
- table.
- partition key.
- document.

Answer :

Answer Area

When using the Azure Cosmos DB Gremlin API, the container resource type is projected as a [graph. ▼]
- graph.
- table.
- partition key.
- document.

Reference:
https://docs.microsoft.com/en-us/azure/cosmos-db/create-graph-gremlin-console

Question 106 (Describe how to work with non-relational data on Azure)

At which two levels can you set the throughput for an Azure Cosmos DB account? Each correct answer presents a complete solution. (Choose two.)
NOTE: Each correct selection is worth one point.

A. database
B. item
C. container
D. partition

Answer : AC

Reference:
https://docs.microsoft.com/en-us/azure/cosmos-db/set-throughput

---

Question 107 (Describe how to work with non-relational data on Azure)

HOTSPOT -
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.
Hot Area:

Answer Area

Statements	Yes	No
Azure Table storage supports multiple read replicas.	○	○
Azure Table storage supports multiple write regions.	○	○
The Azure Cosmos DB Table API supports multiple read replicas.	○	○
The Azure Cosmos DB Table API supports multiple write regions.	○	○

Answer :

Answer Area

Statements	Yes	No
Azure Table storage supports multiple read replicas.	○	○
Azure Table storage supports multiple write regions.	○	○
The Azure Cosmos DB Table API supports multiple read replicas.	○	○
The Azure Cosmos DB Table API supports multiple write regions.	○	○

Reference:
https://docs.microsoft.com/en-us/azure/cosmos-db/high-availability

---

Question 108 (Describe how to work with non-relational data on Azure)

Which setting can only be configured during the creation of an Azure Cosmos DB account?

A. geo-redundancy
B. multi-region writes
C. production or non-production account type
D. API

Answer : D

Explanation:
You must create a separate account for each API.
Reference:
https://docs.microsoft.com/en-us/azure/cosmos-db/how-to-manage-database-account

Question 109 (Describe an analytics workload on Azure)

You need to gather real-time telemetry data from a mobile application.
Which type of workload describes this scenario?

A. Online Transaction Processing (OLTP)
B. batch
C. massively parallel processing (MPP)
D. streaming

Answer : D

Reference:
https://docs.microsoft.com/en-in/azure/azure-monitor/overview

Question 110 (Describe an analytics workload on Azure)

You have a SQL pool in Azure Synapse Analytics that is only used actively every night for eight hours.
You need to minimize the cost of the dedicated SQL pool as much as possible during idle times. The solution must ensure that the data remains intact.
What should you do on the SQL pool?

A. Scale down the data warehouse units (DWUs).
B. Pause the pool.
C. Create a user-defined restore point.
D. Delete the pool

Answer : B

Reference:
https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-manage-compute-overview

Question 111 (Describe an analytics workload on Azure)

Which Azure Data Factory component initiates the execution of a pipeline?

A. a control flow
B. a trigger
C. a parameter
D. an activity

Answer : B

Reference:
https://docs.microsoft.com/en-us/azure/data-factory/concepts-pipeline-execution-triggers#trigger-execution

Question 112 (Describe an analytics workload on Azure)

Your company has a reporting solution that has paginated reports. The reports query a dimensional model in a data warehouse.
Which type of processing does the reporting solution use?

A. stream processing
B. batch processing
C. Online Analytical Processing (OLAP)
D. Online Transaction Processing (OLTP)

Answer : C

Reference:
https://datawarehouseinfo.com/how-does-oltp-differ-from-olap-database/

Question 113 (Describe an analytics workload on Azure)

DRAG DROP -
Match the types of activities to the appropriate Azure Data Factory activities.
To answer, drag the appropriate activity type from the column on the left to its Data Factory activity on the right. Each activity type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.
Select and Place:

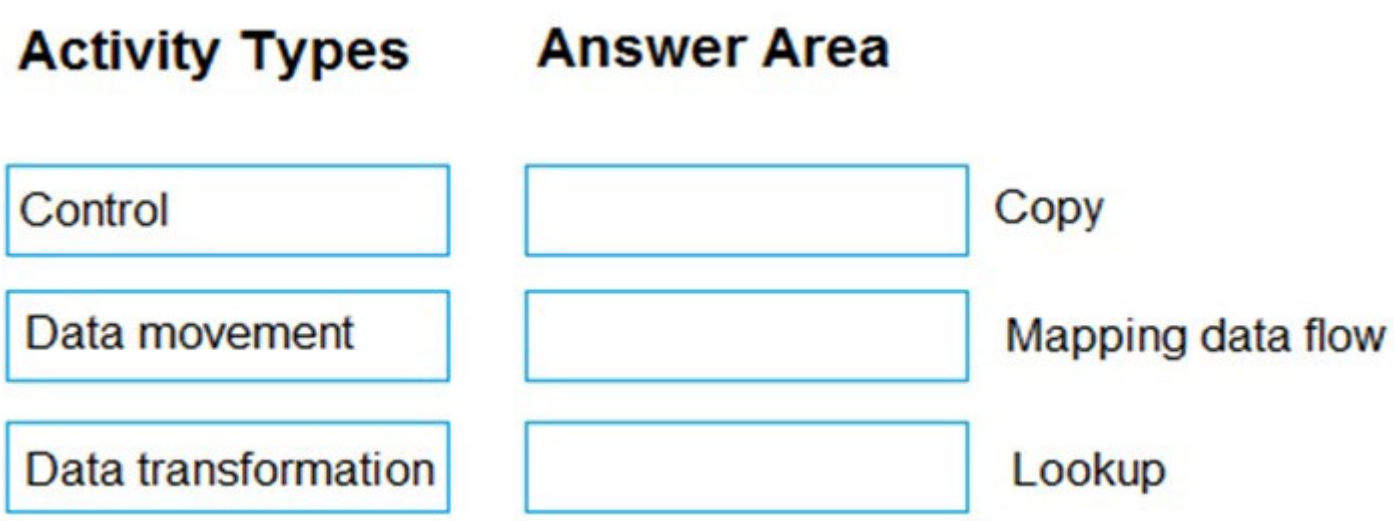

Answer :

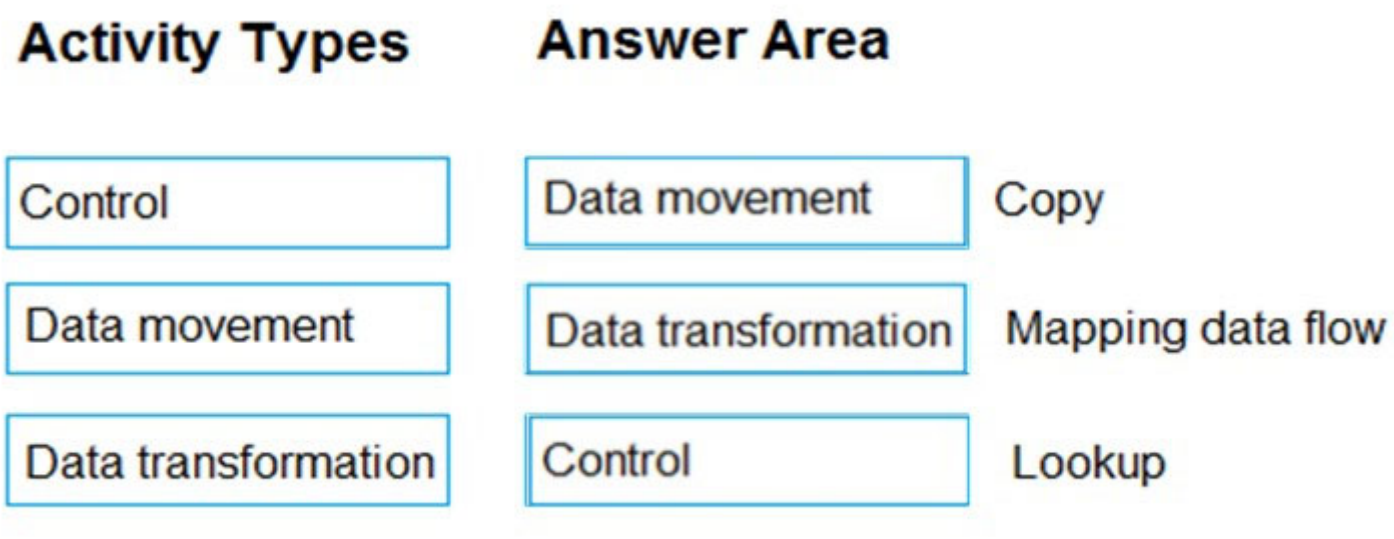

Explanation:

Box 1: Data movement -

Box 2: Data transformation -
A pipeline could contain a set of activities that ingest and clean log data, and then kick off a mapping data flow to analyze the log data.

Box 3: Control -
Lookup Activity is a control flow activity.
Lookup Activity can be used to read or look up a record/ table name/ value from any external source. This output can further be referenced by succeeding activities.
Reference:
https://docs.microsoft.com/en-us/azure/data-factory/concepts-pipelines-activities#data-transformation-activities

Question 114 (Describe an analytics workload on Azure)

What are three characteristics of an Online Transaction Processing (OLTP) workload? Each correct answer presents a complete solution. (Choose three.)
NOTE: Each correct selection is worth one point.

A. denormalized data
B. heavy writes and moderate reads
C. light writes and heavy reads
D. schema on write
E. schema on read
F. normalized data

Answer : BDF

Explanation:
B: Transactional data tends to be heavy writes, moderate reads.
D: Typical traits of transactional data include: schema on write, strongly enforced
F: Transactional data tends to be highly normalized.
Reference:
https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing

Question 115 (Describe an analytics workload on Azure)

Which two activities can be performed entirely by using the Microsoft Power BI service without relying on Power BI Desktop? Each correct answer presents a complete solution. (Choose two.)
NOTE: Each correct selection is worth one point.

A. report and dashboard creation
B. report sharing and distribution
C. data modeling
D. data acquisition and preparation

Answer : AD

Question 116 (Describe an analytics workload on Azure)

HOTSPOT -
To complete the sentence, select the appropriate option in the answer area.
Hot Area:

Answer Area

In Azure Data Factory, you can use [▼] to orchestrate pipeline activities that depend on the output of other pipeline activities.

- a control flow
- a dataset
- a linked service
- an integration runtime

Answer :

Answer Area

In Azure Data Factory, you can use [▼] to orchestrate pipeline activities that depend on the output of other pipeline activities.

- a control flow
- a dataset
- a linked service
- an integration runtime

Reference:
https://docs.microsoft.com/en-us/azure/data-factory/frequently-asked-questions

Question 117 (Describe an analytics workload on Azure)

You need to develop a solution to provide data to executives. The solution must provide an interactive graphical interface, depict various key performance indicators, and support data exploration by using drill down.
What should you use in Microsoft Power BI?

A. a dashboard
B. a report
C. a dataflow
D. Microsoft Power Apps

Answer : B

Reference:
https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards https://docs.microsoft.com/en-us/power-bi/visuals/power-bi-visualization-kpi https://docs.microsoft.com/en-us/power-bi/consumer/end-user-drill

Question 118 (Describe an analytics workload on Azure)

Which two Azure services can be used to provision Apache Spark clusters? Each correct answer presents a complete solution. (Choose two.)
NOTE: Each correct selection is worth one point.

A. Azure Time Series Insights
B. Azure HDInsight
C. Azure Databricks
D. Azure Log Analytics

Answer : BC

Reference:
https://www.sqlshack.com/a-beginners-guide-to-azure-databricks/

Question 119 (Describe an analytics workload on Azure)

You have a quality assurance application that reads data from a data warehouse

You have a quality assurance application that reads data from a data warehouse.
Which type of processing does the application use?

A. Online Transaction Processing (OLTP)
B. batch processing
C. Online Analytical Processing (OLAP)
D. stream processing

Answer : A

Reference:
https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing

---

Question 120 (Describe an analytics workload on Azure)

Which three objects can be added to a Microsoft Power BI dashboard? Each correct answer presents a complete solution. (Choose three.)
NOTE: Each correct selection is worth one point.

A. a report page
B. a Microsoft PowerPoint slide
C. a visualization from a report
D. a dataflow
E. a text box

Answer : ACE

Reference:
https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboards https://docs.microsoft.com/en-us/power-bi/create-reports/service-dashboard-add-widget

---

Question 121 (Describe an analytics workload on Azure)

HOTSPOT -
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.
Hot Area:

Answer Area

Statements	Yes	No
A Microsoft Power BI dashboard is associated with a single workspace.	○	○
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	○	○
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	○	○

Answer :

Answer Area

Statements	Yes	No
A Microsoft Power BI dashboard is associated with a single workspace.	○	○
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	○	○
A Microsoft Power BI dashboard can display visualizations from a Microsoft Excel workbook.	○	○

Reference:
https://docs.microsoft.com/en-us/power-bi/connect-data/service-datasets-across-workspaces https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboardshttps://powerbi.microsoft.com/en-us/excel-and-power-bi/

Question 122 (Describe an analytics workload on Azure)

Which Azure Data Factory component provides the compute environment for activities?

A. SSIS packages
B. an integration runtime
C. a control flow
D. a pipeline

Answer : B

Explanation:
The Integration Runtime (IR) is the compute infrastructure used by Azure Data Factory to provide the following data integration capabilities across different network environments:
⇨ Data Flow
⇨ Data movement
⇨ Activity dispatch
⇨ SSIS package execution
Reference:
https://docs.microsoft.com/en-us/azure/data-factory/concepts-integration-runtime

---

Question 123 (Describe an analytics workload on Azure)

HOTSPOT -
To complete the sentence, select the appropriate option in the answer area.
Hot Area:

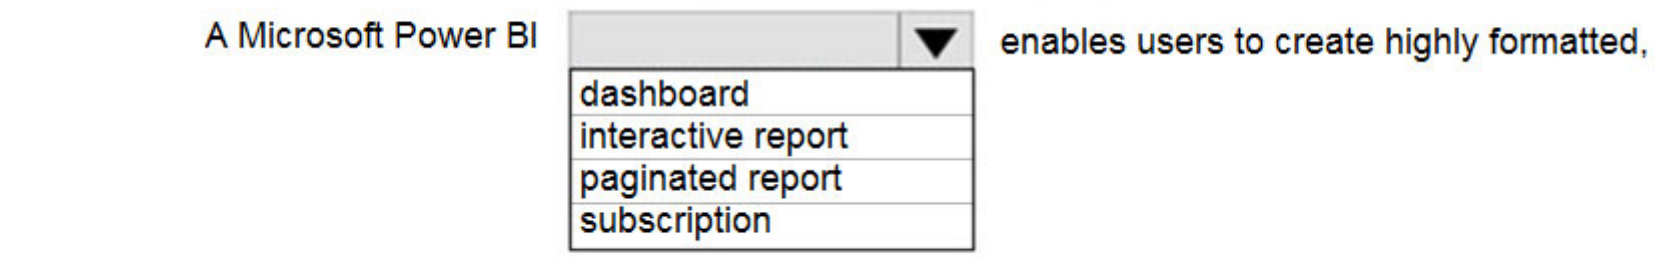

Answer :

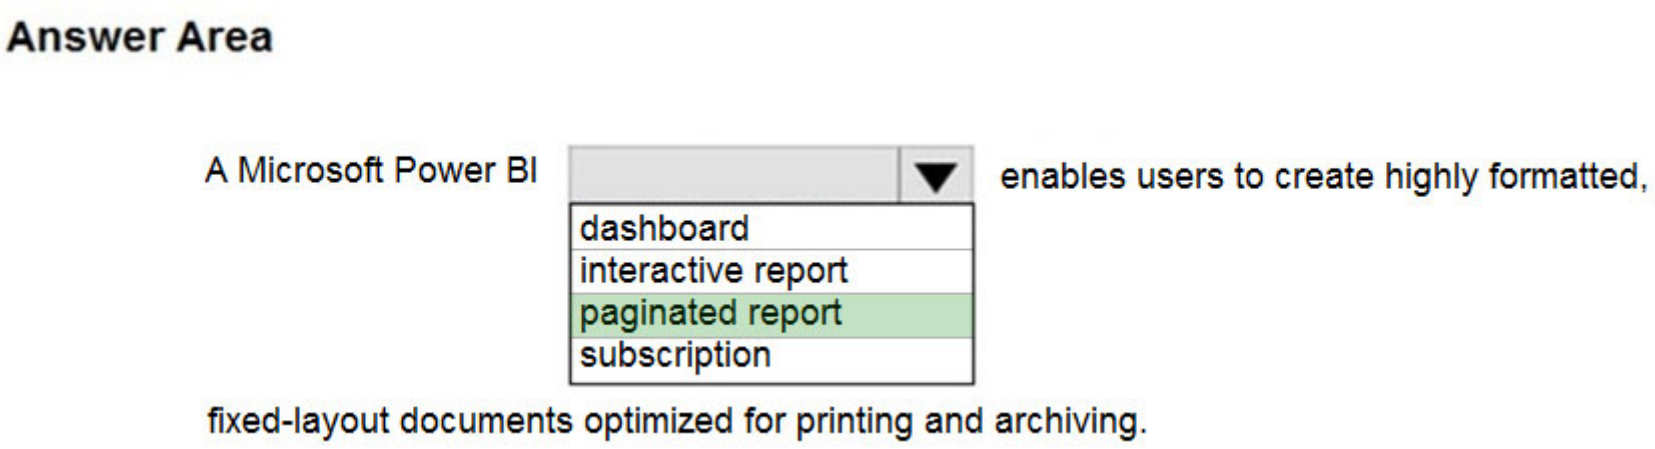

Explanation:
Paginated Reports in Power BI now allows users to generate these fixed-layout documents optimized for printing and archiving, such as PDF and Word files.
These document-style reports with visualizations that provide additional control, like which tables expand horizontally and vertically to display all their data and continue from page to page as needed.
Reference:
https://powerbi.microsoft.com/en-us/blog/announcing-paginated-reports-in-power-bi-general-availability/

Question 124 (Describe an analytics workload on Azure)

What are two uses of data visualization? Each correct answer presents a complete solution.
NOTE: Each correct selection is worth one point.

A. Represent trends and patterns over time
B. Implement machine learning to predict future values
C. Communicate the significance of data
D. Enforce business logic across reports

Answer : AC

Explanation:
Data visualization is a key component in being able to gain insight into your data. It helps make big and small data easier for humans to understand. It also makes it easier to detect patterns, trends, and outliers in groups of data.

Data visualization brings data to help you find key business insights quickly and effectively

Data visualization brings data to help you find key business insights quickly and effectively.
Reference:
https://docs.microsoft.com/en-us/azure/synapse-analytics/spark/apache-spark-data-visualization

Question 125 (Describe an analytics workload on Azure)

You need to use Transact-SQL to query files in Azure Data Lake Storage from an Azure Synapse Analytics data warehouse.
What should you use to query the files?

A. Azure Functions
B. Microsoft SQL Server Integration Services (SSIS)
C. PolyBase
D. Azure Data Factory

Answer : C

Reference:
https://docs.databricks.com/data/data-sources/azure/synapse-analytics.html

Question 126 (Describe an analytics workload on Azure)

HOTSPOT -
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.
Hot Area:

Answer Area

Statements	Yes	No
You can copy a dashboard between Microsoft Power BI workspaces.	○	○
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	○	○
A Microsoft Power BY dashboard can display visualizations from a Microsoft Excel workbook.	○	○

Answer :

Answer Area

Statements	Yes	No
You can copy a dashboard between Microsoft Power BI workspaces.	○	○
A Microsoft Power BI dashboard can only display visualizations from a single dataset.	○	○
A Microsoft Power BY dashboard can display visualizations from a Microsoft Excel workbook.	○	○

Explanation:

Box 1: No -
You can duplicate a dashboard. The duplicate ends up in the same Power BI workspace.
There is no current functionality that allows you to move reports from one workspace to another.

Box 2: No -

Box 3: Yes -
Reference:
https://docs.microsoft.com/en-us/power-bi/connect-data/service-datasets-across-workspaces https://docs.microsoft.com/en-us/power-bi/consumer/end-user-dashboardshttps://powerbi.microsoft.com/en-us/excel-and-power-bi/

Question 127 (Describe an analytics workload on Azure)

What are three characteristics of an Online Transaction Processing (OLTP) workload? Each correct answer presents a complete solution. (Choose three.)
NOTE: Each correct selection is worth one point.

A. denormalized data
B. heavy writes and moderate reads
C. light writes and heavy reads
D. schema defined in a database
E. schema defined when reading unstructured data from a database
F. normalized data

Answer : BDF

Explanation:
B: Transactional data tends to be heavy writes, moderate reads.
D: Typical traits of transactional data include: schema on write, strongly enforced. The schema is defined in a database.
F: Transactional data tends to be highly normalized.
Reference:
https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/online-transaction-processing

Question 128 (Describe an analytics workload on Azure)

What is the primary purpose of a data warehouse?

A. to provide answers to complex queries that rely on data from multiple sources
B. to provide transformation services between source and target data stores
C. to provide read-only storage of relational and non-relational historical data
D. to provide storage for transactional line-of-business (LOB) applications

Answer : C

Explanation:
Consider using a data warehouse when you need to keep historical data separate from the source transaction systems for performance reasons. Data warehouses make it easy to access historical data from multiple locations, by providing a centralized location using common formats, keys, and data models.
Query both relational and nonrelational data.
Incorrect Answers:
D: Data warehouses don't need to follow the same terse data structure you may be using in your OLTP databases.
Reference:
https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/data-warehousing

Question 129 (Describe an analytics workload on Azure)

DRAG DROP -
Match the Azure services to the appropriate locations in the architecture.
To answer, drag the appropriate service from the column on the left to its location on the right. Each service may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.
Select and Place:

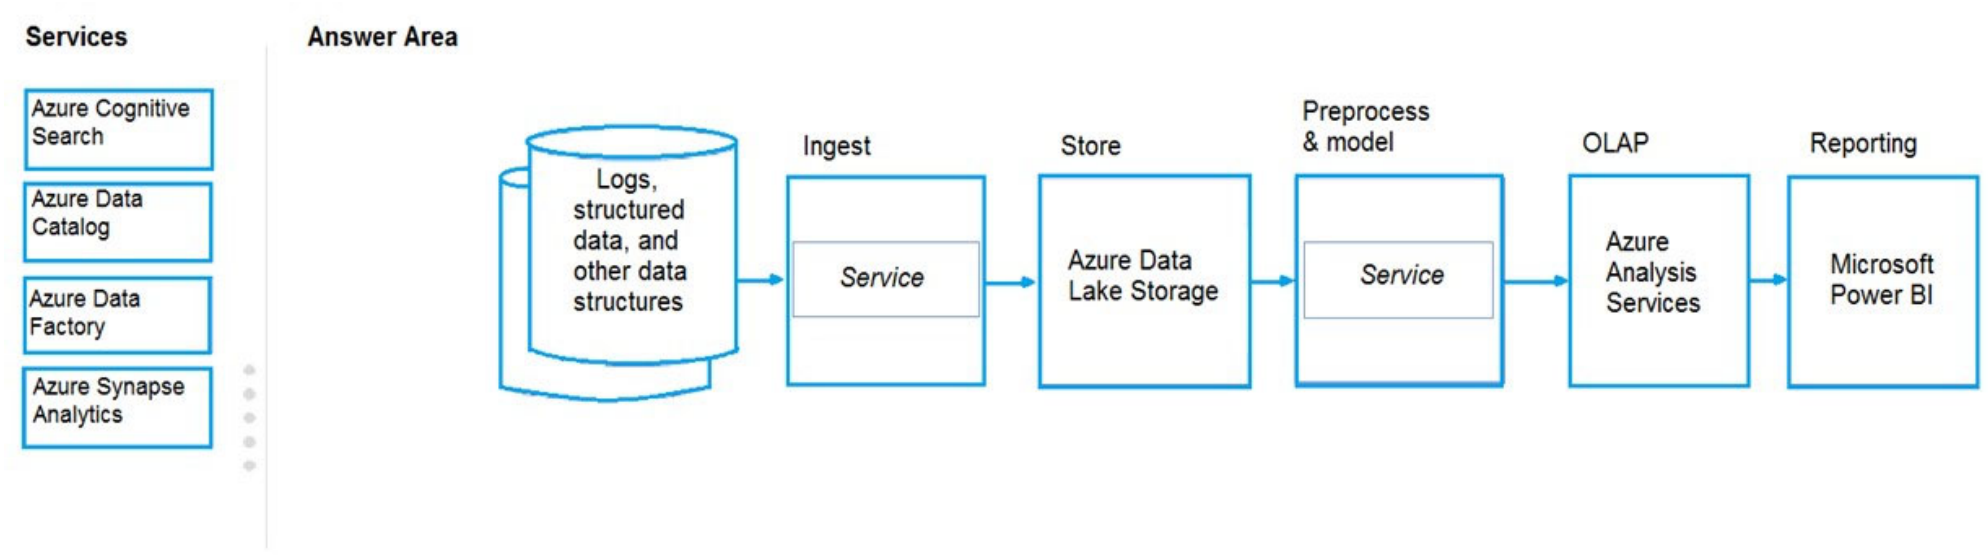

Answer :

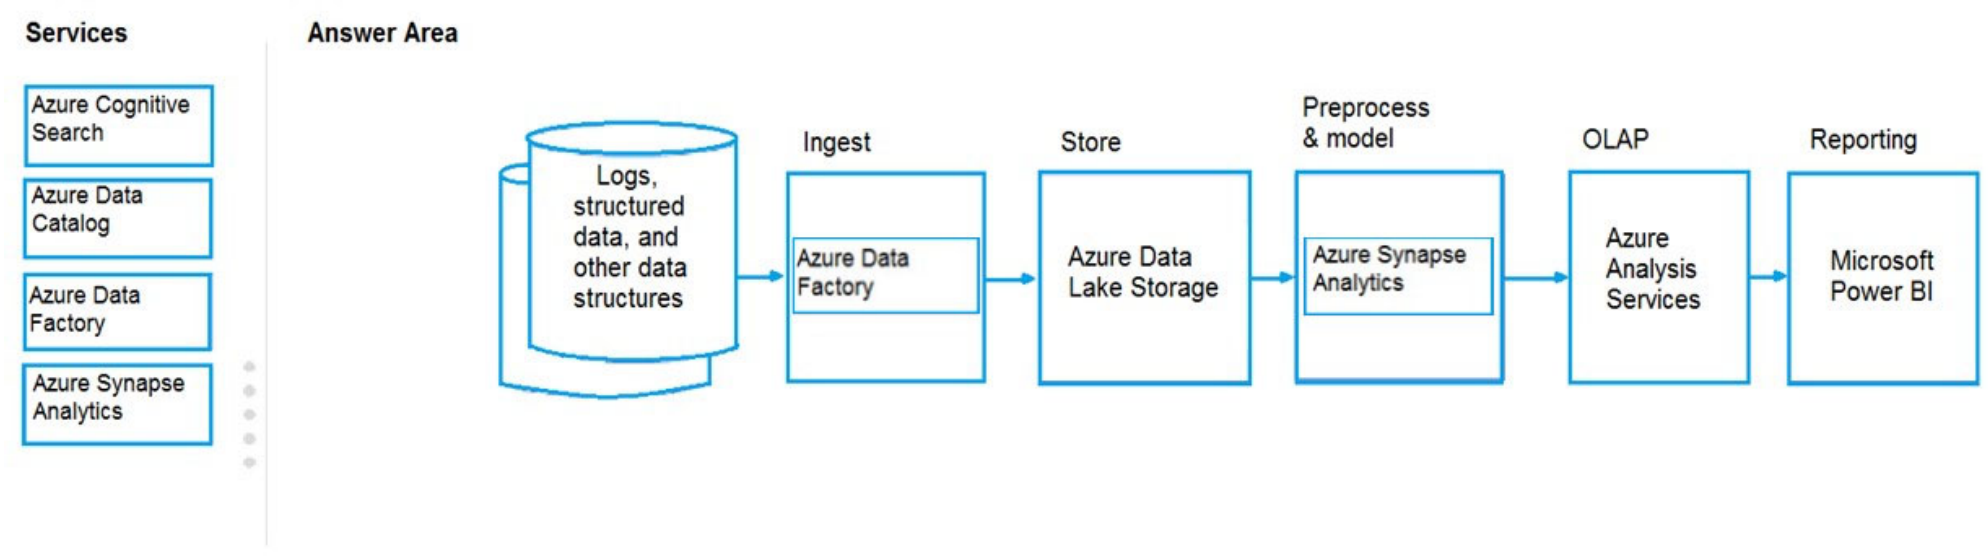

Explanation:

Box Ingest: Azure Data Factory -
You can build a data ingestion pipeline with Azure Data Factory (ADF).
Box Preprocess & model: Azure Synapse Analytics
Use Azure Synapse Analytics to preprocess data and deploy machine learning models.
Reference:
https://docs.microsoft.com/en-us/azure/machine-learning/how-to-data-ingest-adf https://docs.microsoft.com/en-us/azure/machine-learning/team-data-science-process/sqldw-walkthrough

Question 130 (Describe an analytics workload on Azure)

DRAG DROP -
Match the types of workloads to the appropriate scenarios.
To answer, drag the appropriate workload type from the column on the left to its scenario on the right. Each workload type may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.
Select and Place:

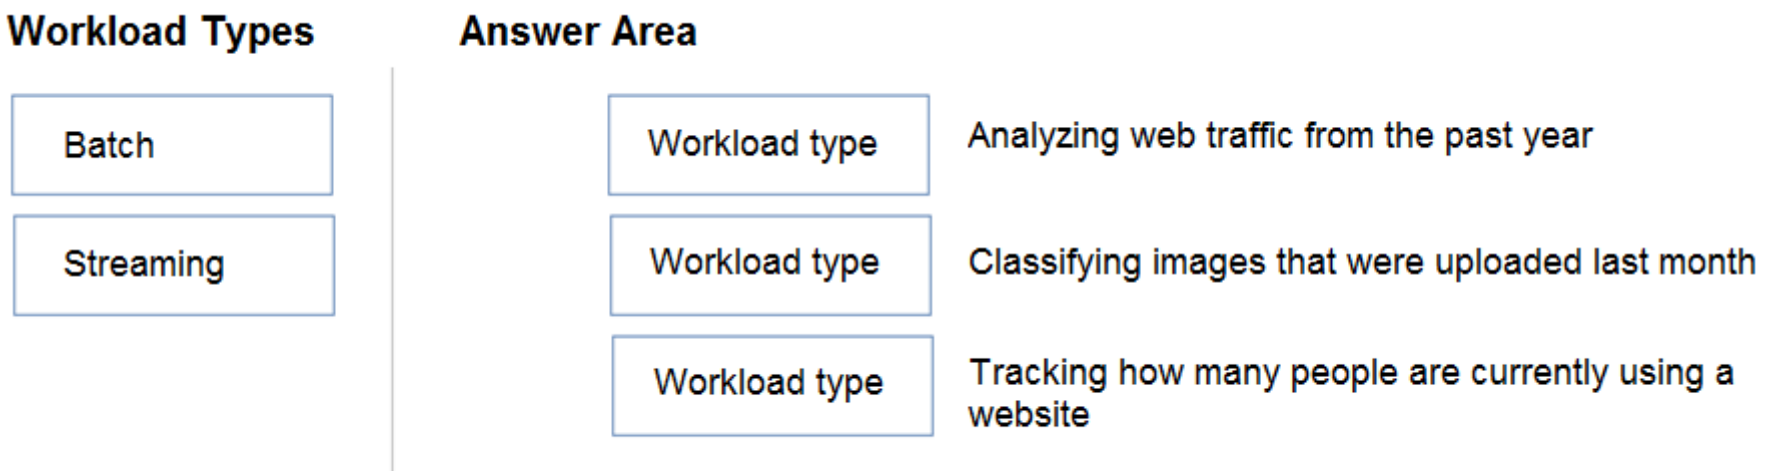

Answer :

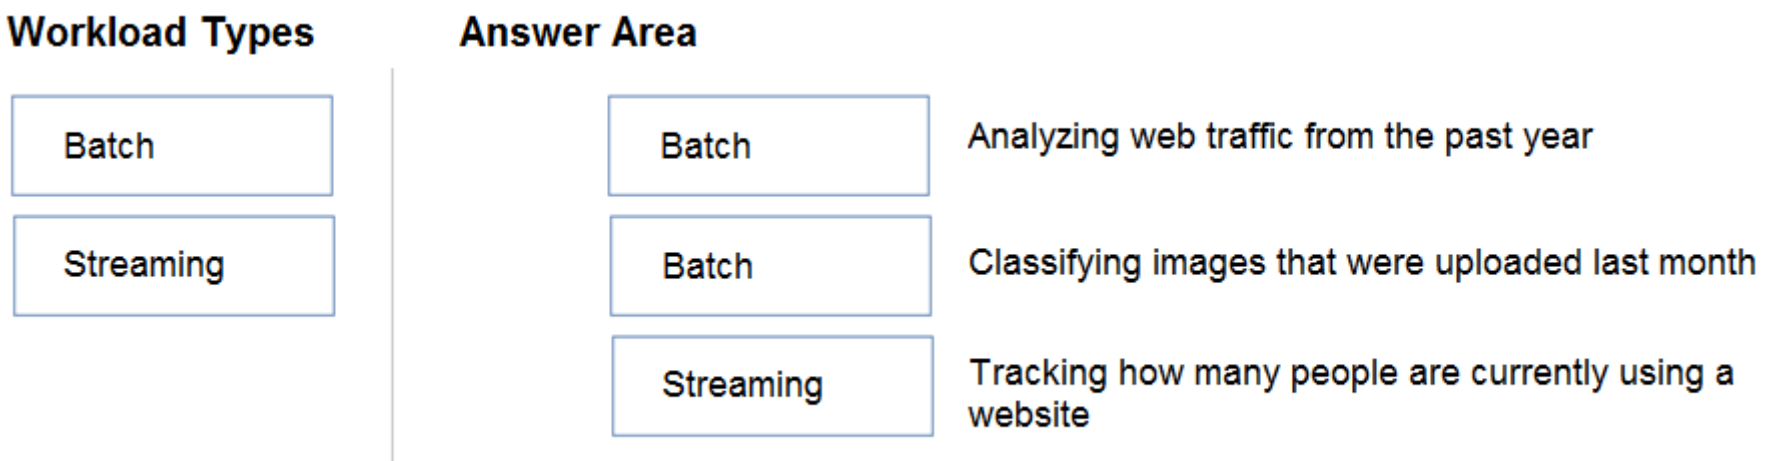

Explanation:

Box 1: Batch -
The batch processing model requires a set of data that is collected over time while the stream processing model requires data to be fed into an analytics tool, often in micro-batches, and in real-time.
The batch Processing model handles a large batch of data while the Stream processing model handles individual records or micro-batches of few records.
In Batch Processing, it processes over all or most of the data but in Stream Processing, it processes over data on a rolling window or most recent record.

Box 2: Batch -

Box 3: Streaming -
Reference:
https://k21academy.com/microsoft-azure/dp-200/batch-processing-vs-stream-processing

Question 131 (Describe an analytics workload on Azure)

HOTSPOT -
To complete the sentence, select the appropriate option in the answer area.
Hot Area:

Answer Area

In a data warehousing workload, data

- from a single source is distributed to multiple locations
- from multiple sources is combined in a single location
- is added to a queue for multiple systems to process
- is used to train machine learning models

Answer :

Answer Area

In a data warehousing workload, data

- from a single source is distributed to multiple locations
- from multiple sources is combined in a single location
- is added to a queue for multiple systems to process
- is used to train machine learning models

Explanation:
Note: The data warehouse workload encompasses:
➪ The entire process of loading data into the warehouse
➪ Performing data warehouse analysis and reporting
➪ Managing data in the data warehouse
➪ Exporting data from the data warehouse
Reference:
https://docs.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/sql-data-warehouse-workload-management

Question 132 (Describe an analytics workload on Azure)

HOTSPOT -
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.
Hot Area:

Answer Area

Statements	Yes	No
A pipeline is a representation of a data structure within Azure Data Factory	○	○
Azure Data Factory pipelines can execute other pipelines	○	○
A processing step within an Azure Data Factory pipeline is an activity	○	○

Answer :

Answer Area

Statements	Yes	No
A pipeline is a representation of a data structure within Azure Data Factory	○	○
Azure Data Factory pipelines can execute other pipelines	○	○
A processing step within an Azure Data Factory pipeline is an activity	○	○

Explanation:

Explanation:

Box 1: No -
A pipeline is a logical grouping of activities that together perform a task.

Box 2: Yes -
You can construct pipeline hierarchies with data factory.

Box 3: Yes -
A pipeline is a logical grouping of activities that together perform a task.
Reference:
https://mrpaulandrew.com/2019/09/25/azure-data-factory-pipeline-hierarchies-generation-control/

Question 133 (Describe an analytics workload on Azure)

DRAG DROP -
Match the Azure services to the appropriate requirements.
To answer, drag the appropriate service from the column on the left to its requirement on the right. Each service may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.
Select and Place:

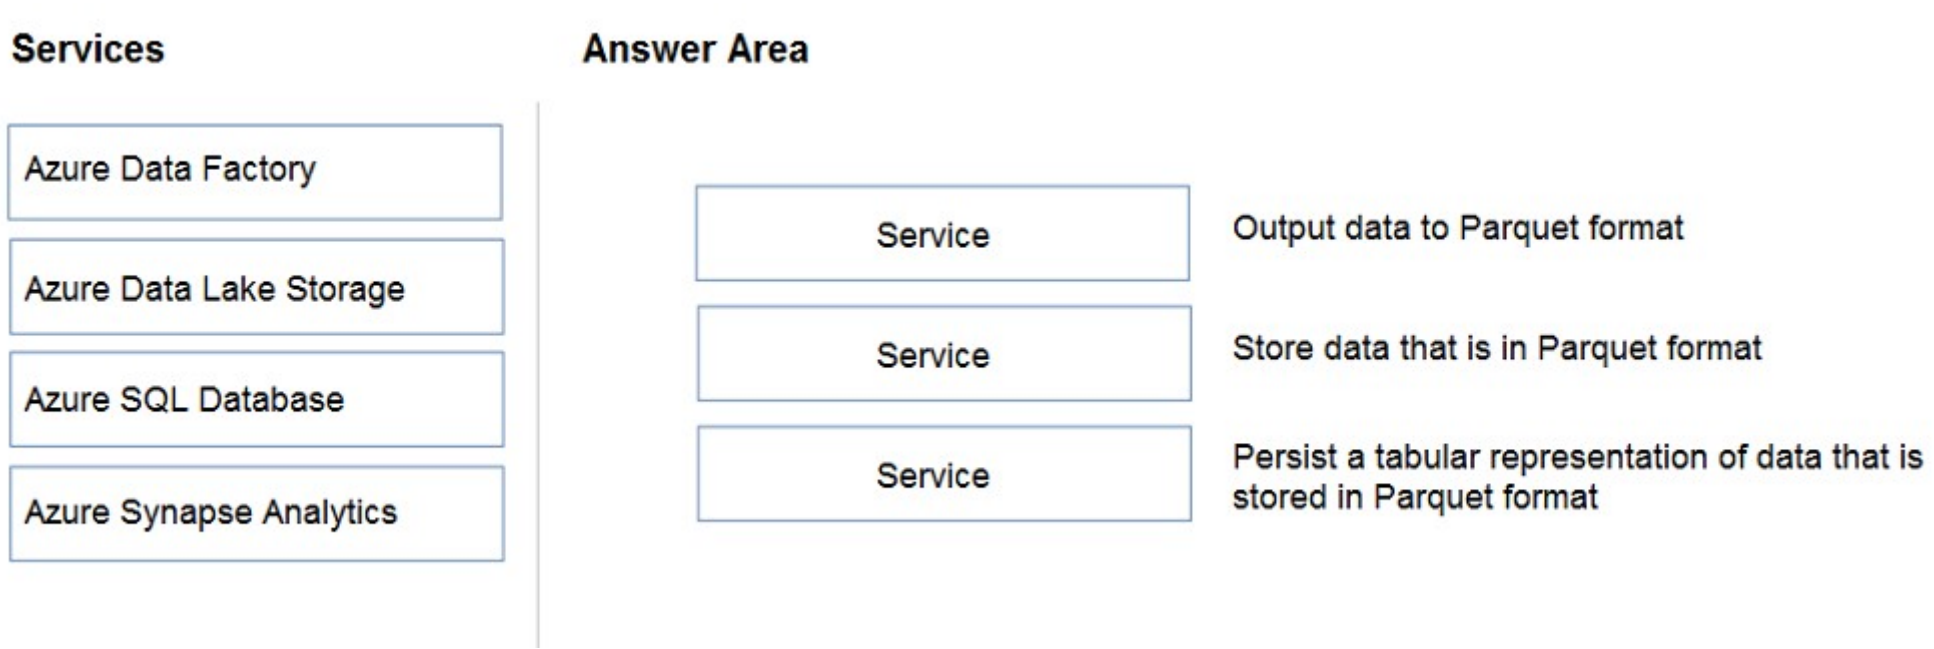

Answer :

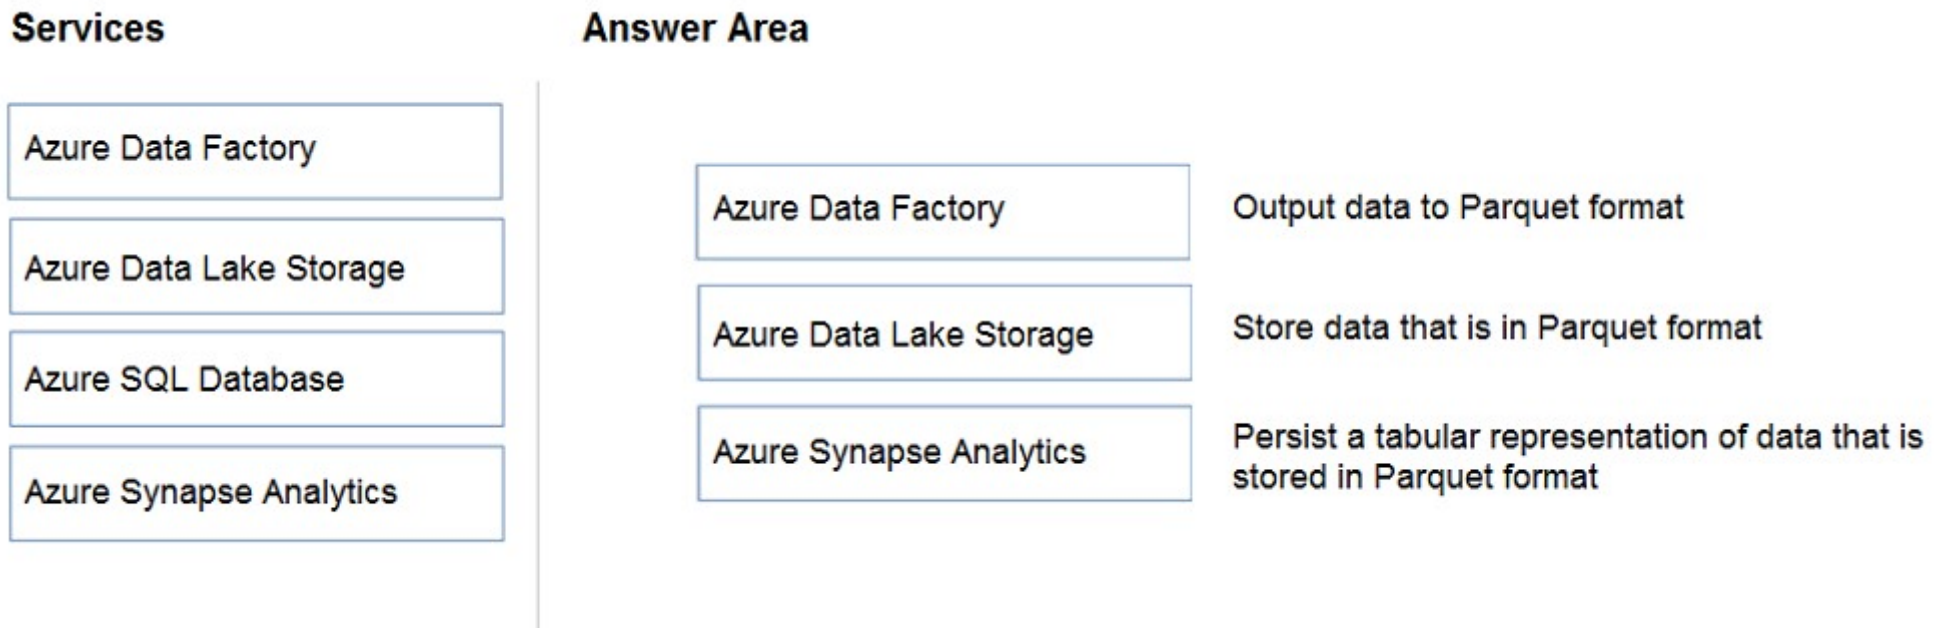

Explanation:

Box 1: Azure Data Factory -

Box 2: Azure Data Lake Storage -
Azure Data Lake Storage (ADLA) now natively supports Parquet files. ADLA adds a public preview of the native extractor and outputter for the popular Parquet file format

Box 3: Azure Synapse Analytics -
Use Azure Synapse Analytics Workspaces.
Reference:
https://docs.microsoft.com/en-us/azure/data-factory/supported-file-formats-and-compression-codecs

---

Question 134 (Describe an analytics workload on Azure)

HOTSPOT -
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.
Hot Area:

Answer Area

Statements	Yes	No
Azure Synapse Analytics scales storage and compute independently	○	○
Azure Synapse Analytics can be paused to reduce compute costs	○	○
An Azure Synapse Analytics data warehouse has a fixed storage capacity	○	○

Answer :

Answer Area

Statements	Yes	No
Azure Synapse Analytics scales storage and compute independently	○	○
Azure Synapse Analytics can be paused to reduce compute costs	○	○
An Azure Synapse Analytics data warehouse has a fixed storage capacity	○	○

Explanation:

Box 1: Yes -
Compute is separate from storage, which enables you to scale compute independently of the data in your system.

Box 2: Yes -
You can use the Azure portal to pause and resume the dedicated SQL pool compute resources.
Pausing the data warehouse pauses compute. If your data warehouse was paused for the entire hour, you will not be charged compute during that hour.

Box 3: No -
Storage is sold in 1 TB allocations. If you grow beyond 1 TB of storage, your storage account will automatically grow to 2 TBs.
Reference:
https://azure.microsoft.com/en-us/pricing/details/synapse-analytics/

Question 135 (Describe an analytics workload on Azure)

What should you use to build a Microsoft Power BI paginated report?

A. Charticulator
B. Power BI Desktop
C. the Power BI service
D. Power BI Report Builder

Answer : D

Explanation:
Power BI Report Builder is the standalone tool for authoring paginated reports for the Power BI service.
Reference:
https://docs.microsoft.com/en-us/power-bi/paginated-reports/paginated-reports-report-builder-power-bi

---

Question 136 (Describe an analytics workload on Azure)

DRAG DROP -
Match the Azure services to the appropriate locations in the architecture.
To answer, drag the appropriate service from the column on the left to its location on the right. Each service may be used once, more than once, or not at all.
NOTE: Each correct match is worth one point.
Select and Place:

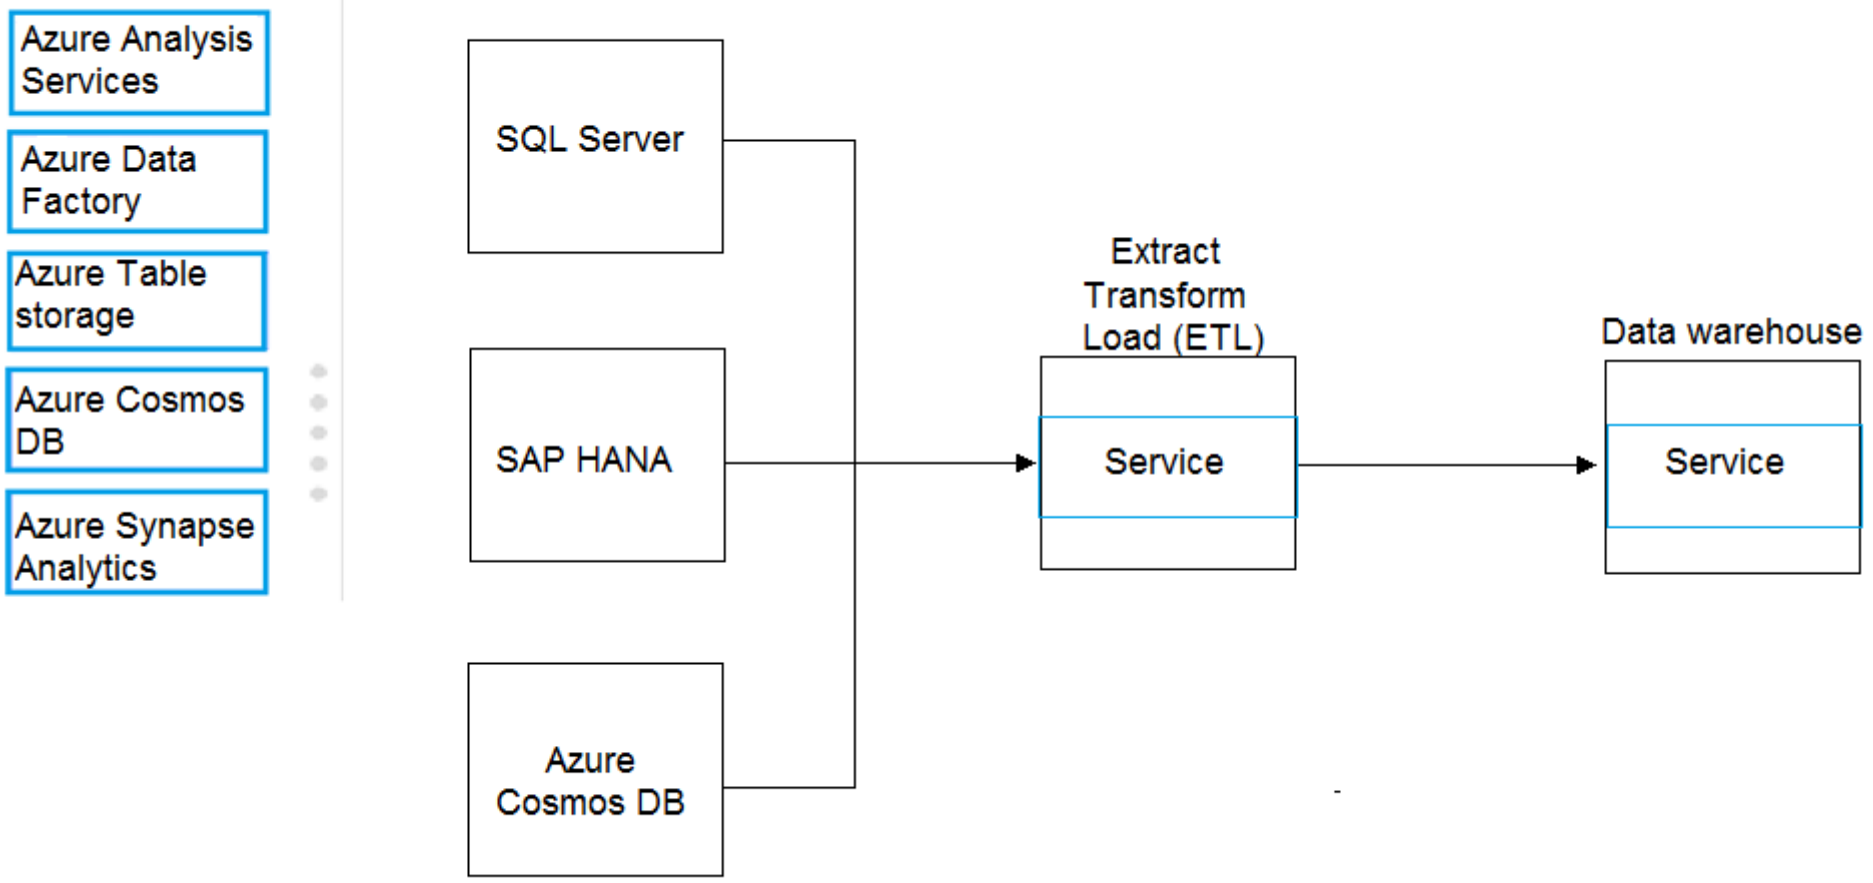

Answer :

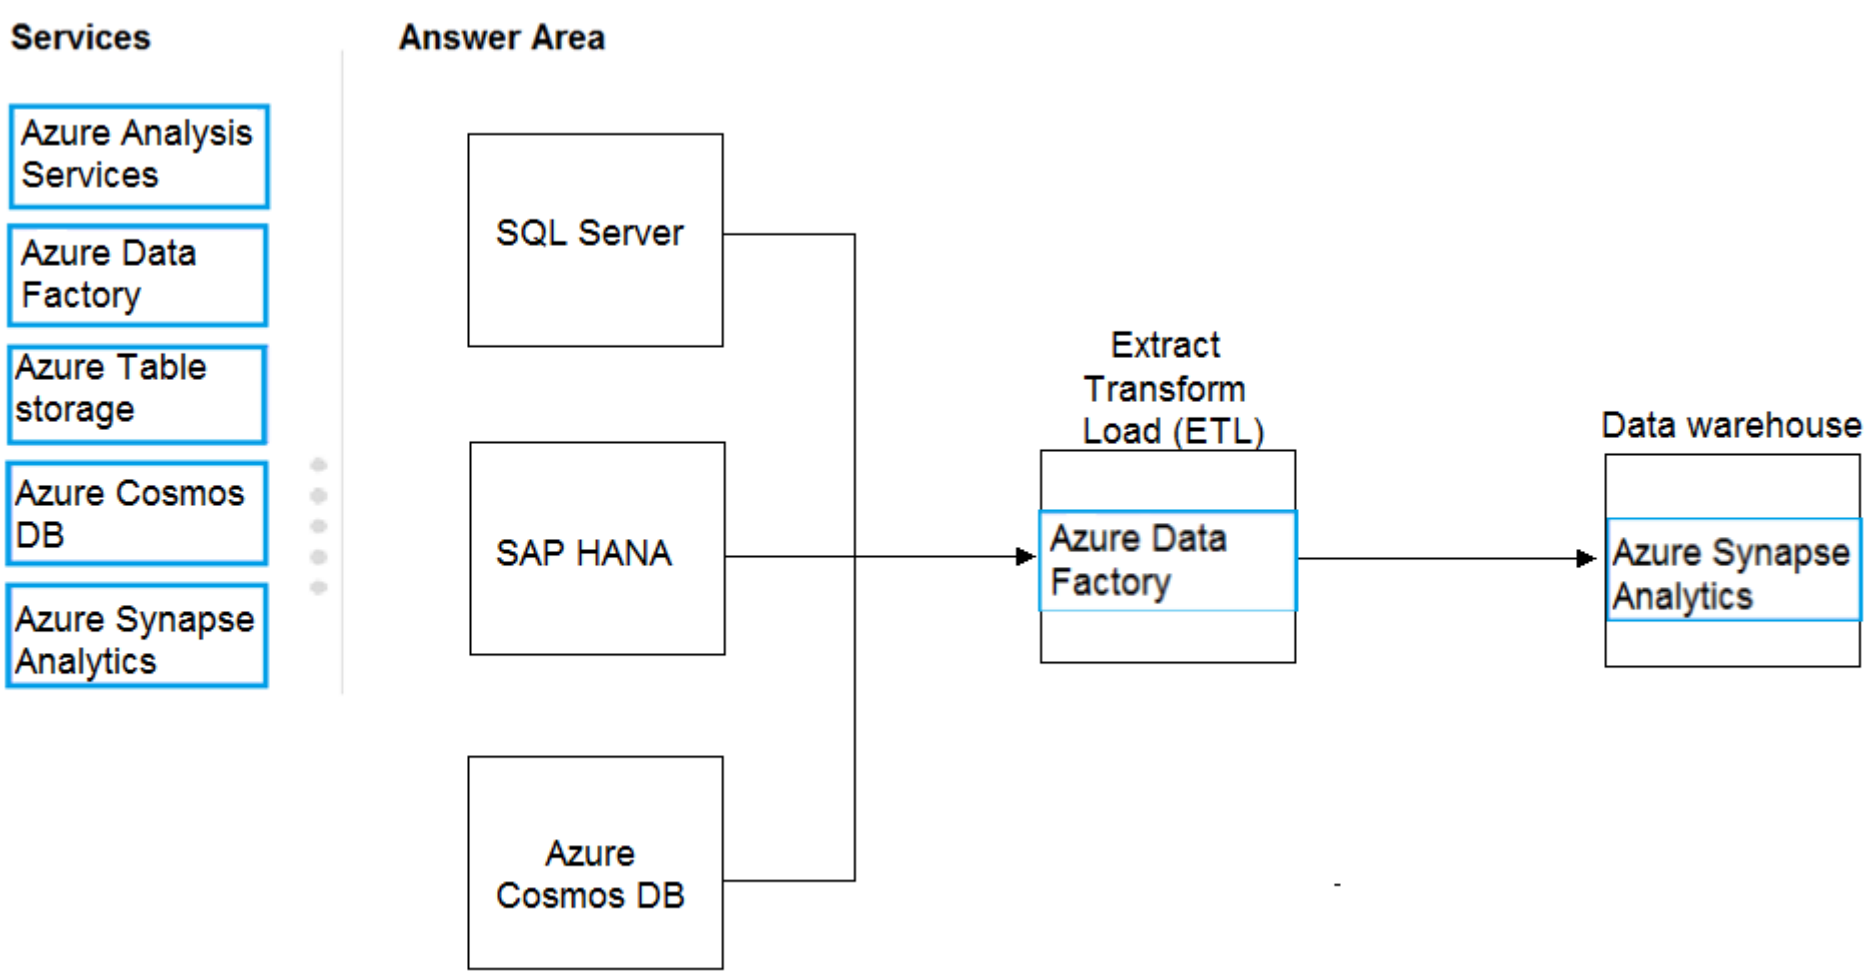

Explanation:

Box 1: Azure Data factory -
Relevant Azure service for the three ETL phases are Azure Data Factory and SQL Server Integration Services (SSIS).

Box 2: Azure Synapse Analytics -
You can copy and transform data in Azure Synapse Analytics by using Azure Data Factory
Note: Azure Synapse Analytics connector is supported for the following activities:
- Copy activity with supported source/sink matrix table
- Mapping data flow
- Lookup activity
- GetMetadata activity

Reference:
https://docs.microsoft.com/en-us/azure/architecture/data-guide/relational-data/etl https://docs.microsoft.com/en-us/azure/data-factory/connector-azure-sql-data-warehouse

---

Question 137 (Describe an analytics workload on Azure)

HOTSPOT -

For each of the following statements, select Yes if the statement is true. Otherwise, select No

For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.
Hot Area:

Answer Area

Statements	Yes	No
Azure Databricks is an Apache Spark-based collaborative analytics platform.	○	○
Azure Analysis Services is used for transactional workloads.	○	○
Azure Data Factory orchestrates data integration workflows.	○	○

Answer :

Answer Area

Statements	Yes	No
Azure Databricks is an Apache Spark-based collaborative analytics platform.	○	○
Azure Analysis Services is used for transactional workloads.	○	○
Azure Data Factory orchestrates data integration workflows.	○	○

Reference:
https://docs.microsoft.com/en-us/azure/databricks/scenarios/what-is-azure-databricks https://docs.microsoft.com/en-us/azure/analysis-services/analysis-services-overview https://docs.microsoft.com/en-us/azure/data-factory/introduction

Question 138 (Describe an analytics workload on Azure)

Which scenario is an example of a streaming workload?

A. sending transactions that are older than a month to an archive
B. sending transactions daily from point of sale (POS) devices
C. sending telemetry data from edge devices
D. sending cloud infrastructure metadata every 30 minutes

Answer : C

Question 139 (Describe an analytics workload on Azure)

HOTSPOT -
To complete the sentence, select the appropriate option in the answer area.
Hot Area:

Answer Area

Batch workloads [▼]

- process data in memory, row-by-row.
- collect and process data at most once a day.
- process data as new data is received in near real-time.
- collect data and then process the data when a condition is met.

Answer :

Answer Area

Batch workloads

process data in memory, row-by-row.
collect and process data at most once a day.
process data as new data is received in near real-time

Question 140 (Describe an analytics workload on Azure)

HOTSPOT -
For each of the following statements, select Yes if the statement is true. Otherwise, select No.
NOTE: Each correct selection is worth one point.
Hot Area:

Answer Area

Statements	Yes	No
Processing salary payments once a month is an example of a batch workload.	○	○
A wind turbine that sends 50 sensor readings per second is an example of a streaming workload.	○	○
A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload.	○	○

Answer :

Answer Area

Statements	Yes	No
Processing salary payments once a month is an example of a batch workload.	○	○
A wind turbine that sends 50 sensor readings per second is an example of a streaming workload.	○	○
A home electricity meter that sends readings once a day to an energy provider is an example of a streaming workload.	○	○

Question 141 (Describe an analytics workload on Azure)

A bar chart showing year-to-date sales by region is an example of which type of analytics?

A. predictive
B. prescriptive
C. descriptive
D. diagnostic

Answer : D